

## Implementation file

stock-price-prediction/ |—— data/ |   └—— AAPL.csv ← downloaded dataset |—— model/ |  
                                  └—— lstm\_stock\_model.h5 ← saved trained model |—— stock\_predictor.ipynb ← Jupyter  
notebook for training   └—— app.py ← Streamlit web app (optional)

```
import pandas as pd
```

```
# Load your CSV file
```

```
df = pd.read_csv("Apple Dataset.csv") # Update filename as needed
```

```
print(df.head())
```

```
print(df.columns)
```

```
↔
```

|   | Date       | Open     | High     | Low      | Close    | Adj Close | Volume    |
|---|------------|----------|----------|----------|----------|-----------|-----------|
| 0 | 1980-12-12 | 0.128348 | 0.128906 | 0.128348 | 0.128348 | 0.099058  | 469033600 |
| 1 | 1980-12-15 | 0.122210 | 0.122210 | 0.121652 | 0.121652 | 0.093890  | 175884800 |
| 2 | 1980-12-16 | 0.113281 | 0.113281 | 0.112723 | 0.112723 | 0.086999  | 105728000 |
| 3 | 1980-12-17 | 0.115513 | 0.116071 | 0.115513 | 0.115513 | 0.089152  | 86441600  |
| 4 | 1980-12-18 | 0.118862 | 0.119420 | 0.118862 | 0.118862 | 0.091737  | 73449600  |

```
Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume'], dtype=''
```

## Preprocessing & Normalization

```
import numpy as np
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
# Use only the 'Close' column
```

```
data = df[['Close']].values
```

```
# Normalize
```

```
scaler = MinMaxScaler()
```

```
scaled_data = scaler.fit_transform(data)
```

```
# Create sequences
```

```
def create_sequences(data, sequence_length):
```

```
    X, y = [], []
```

```
    for i in range(sequence_length, len(data)):
```

```
        X.append(data[i-sequence_length:i])
```

```
        y.append(data[i])
```

```
    return np.array(X), np.array(y)
```

```
sequence_length = 60
```

```
X, y = create_sequences(scaled_data, sequence_length)
```

```
# Train-test split
```

```
split = int(0.8 * len(X))
```

```
X_train, X_test = X[:split], X[split:]
```

```
y_train, y_test = y[:split], y[split:]
```

```
# Reshape for LSTM
```

```
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
```

```
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))
```

## Build and Train LSTM Model

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout

model = Sequential([
    LSTM(50, return_sequences=True, input_shape=(X_train.shape[1], 1)),
    Dropout(0.2),
    LSTM(50),
    Dropout(0.2),
    Dense(1)
])

model.compile(optimizer='adam', loss='mean_squared_error')
model.fit(X_train, y_train, epochs=20, batch_size=32, validation_data=(X_test, y_test))

```



```

Epoch 1/20
/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/rnn.py:200: UserWarning:
    super().__init__(**kwargs)
273/273 ————— 11s 34ms/step - loss: 1.4888e-04 - val_loss:
Epoch 2/20
273/273 ————— 9s 32ms/step - loss: 2.0405e-05 - val_loss: 0
Epoch 3/20
273/273 ————— 9s 33ms/step - loss: 1.4847e-05 - val_loss: 6
Epoch 4/20
273/273 ————— 9s 34ms/step - loss: 1.5058e-05 - val_loss: 0
Epoch 5/20
273/273 ————— 9s 30ms/step - loss: 1.5633e-05 - val_loss: 0
Epoch 6/20
273/273 ————— 10s 30ms/step - loss: 1.5764e-05 - val_loss:
Epoch 7/20
273/273 ————— 11s 32ms/step - loss: 1.4503e-05 - val_loss:
Epoch 8/20
273/273 ————— 10s 32ms/step - loss: 1.3796e-05 - val_loss:
Epoch 9/20
273/273 ————— 10s 38ms/step - loss: 1.4618e-05 - val_loss:
Epoch 10/20
273/273 ————— 9s 32ms/step - loss: 1.8402e-05 - val_loss: 0
Epoch 11/20
273/273 ————— 9s 32ms/step - loss: 1.4237e-05 - val_loss: 0
Epoch 12/20
273/273 ————— 10s 33ms/step - loss: 1.3048e-05 - val_loss:
Epoch 13/20
273/273 ————— 8s 31ms/step - loss: 1.3743e-05 - val_loss: 6
Epoch 14/20
273/273 ————— 11s 32ms/step - loss: 1.3134e-05 - val_loss:
Epoch 15/20
273/273 ————— 10s 32ms/step - loss: 1.4271e-05 - val_loss:
Epoch 16/20
273/273 ————— 10s 32ms/step - loss: 1.3002e-05 - val_loss:
Epoch 17/20
273/273 ————— 9s 32ms/step - loss: 1.3636e-05 - val_loss: 0
Epoch 18/20
273/273 ————— 10s 32ms/step - loss: 1.2479e-05 - val_loss:
Epoch 19/20
273/273 ————— 10s 32ms/step - loss: 1.3556e-05 - val_loss:
Epoch 20/20
273/273 ————— 8s 30ms/step - loss: 1.2083e-05 - val_loss: 0
<keras.src.callbacks.history.History at 0x7a15d641ded0>

```

Evaluate and Visualize Predictions

```

import matplotlib.pyplot as plt

predicted = model.predict(X_test)
predicted = scaler.inverse_transform(predicted)
y_test_scaled = scaler.inverse_transform(y_test.reshape(-1, 1))


plt.figure(figsize=(14,6))
plt.plot(y_test_scaled, label='Actual Price')
plt.plot(predicted, label='Predicted Price')
plt.title('Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel('Price')
plt.legend()
plt.show()

```



69/69

1s 11ms/step



```

model.save('lstm_stock_model.keras')

```

