

## Phase-3

**Student Name:** Sasivardhan S K

**Register Number:** 410723106029

**Institution:** Dhanalakshmi College of Engineering

**Department:** Electronics and Communication Engineering

**Date of Submission:** 12/05/2025

**Github Repository Link:**

[https://github.com/SKSasivardhan/NM\\_sasivardhan--DS](https://github.com/SKSasivardhan/NM_sasivardhan--DS)

---

### Cracking the Market Code with AI-Driven Stock Price Prediction Using Time Series Analysis

#### 1. Problem Statement

Stock market price prediction is a critical yet challenging financial task, given the complex and volatile nature of market movements. Accurate stock price forecasting can provide significant business value, aiding traders, investors, and financial institutions in making informed decisions. This project aims to leverage AI and time series analysis to predict future stock prices based on historical data. The goal is to identify patterns and trends that can be used to anticipate market behavior, potentially increasing profit margins and reducing financial risks. This is a regression problem, as the target output is continuous stock prices over time.

## 2. Abstract

This project explores the application of AI-driven time series analysis for stock price prediction. The primary objective is to develop a model capable of accurately forecasting stock prices based on historical data, leveraging machine learning algorithms like ARIMA, LSTM, or Prophet. The approach includes data preprocessing, feature engineering, exploratory data analysis, and rigorous model evaluation to ensure high prediction accuracy. The anticipated outcome is a robust predictive model that can assist in financial decision-making, offering insights into market trends and investment strategies.

## 3. System Requirements

**Hardware:** Minimum 8GB RAM, Intel i5 or higher processor, 256GB SSD.

**Operating System:** Windows 10/11, Linux (Ubuntu 20.04+), or macOS.

**Software:** Python 3.9+, Jupyter Notebook, Colab (optional), required libraries (pandas, numpy, matplotlib, seaborn, scikit-learn, statsmodels, tensorflow, keras)

**Libraries:** NumPy, Pandas, Scikit-learn, Matplotlib, Seaborn.


**Dataset:** Kaggle or bank-provided credit card transaction data (CSV format).

**Optional Tools:** Flask for deployment, Git for version control.

## 4. Objectives

1. To accurately predict future stock prices using historical market data.
2. To identify key features and trends influencing stock price movements.
3. To compare multiple time series forecasting models for performance.
4. To develop a deployable AI solution for real-time market forecasting.

## 5. Flowchart of Project Workflow



## 6. Dataset Description

**Link:** <https://www.kaggle.com/datasets/mrsimple07/stock-price-prediction>

**Source:** Publicly available dataset from Kaggle –Stock price prediction.

**Size:** Contains 336 company with 6 stock (highly imbalanced).

**Features:** 6 stocks and different company stock prices with quotation.

**Label:** Binary class – 0 for legitimate and 1 for fraudulent transactions.

**Format:** CSV file, suitable for supervised machine learning tasks.

**Use:** Used for training, testing, and validating fraud detection models.

**Type:** Public

```

Untitled4.ipynb
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text
Connect
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

[ ] data=pd.read_csv("../stock_data.csv")

[ ] data

    Unnamed: 0  Stock_1  Stock_2  Stock_3  Stock_4  Stock_5
0  2020-01-01  101.764052  100.160928  99.494642  99.909756  101.761266
1  2020-01-02  102.171269  99.969968  98.682973  100.640755  102.528643
2  2020-01-03  103.171258  99.575237  98.182139  100.574847  101.887811
3  2020-01-04  105.483215  99.308641  97.149381  100.925017  101.490049
4  2020-01-05  107.453175  98.188428  99.575396  101.594411  101.604283
...
360 2020-12-26  92.684784  63.408103  98.288992  117.788079  102.995720
361 2020-12-27  92.688279  62.816639  98.061845  116.605106  102.718260
362 2020-12-28  93.551993  63.597651  96.454800  115.441164  103.566068
363 2020-12-29  93.870037  64.114492  95.747485  113.856107  103.257107
364 2020-12-30  93.855317  64.491011  97.805648  112.640418  102.304226
365 rows x 6 columns

[ ] data.head()

    Unnamed: 0  Stock_1  Stock_2  Stock_3  Stock_4  Stock_5
0  2020-01-01  101.764052  100.160928  99.494642  99.909756  101.761266
1  2020-01-02  102.171269  99.969968  98.682973  100.640755  102.528643
2  2020-01-03  103.171258  99.575237  98.182139  100.574847  101.887811
3  2020-01-04  105.483215  99.308641  97.149381  100.925017  101.490049
4  2020-01-05  107.453175  98.188428  99.575396  101.594411  101.604283

[ ] data.drop_duplicates(inplace=True)

[ ] data

    Unnamed: 0  Stock_1  Stock_2  Stock_3  Stock_4  Stock_5
0  2020-01-01  101.764052  100.160928  99.494642  99.909756  101.761266
1  2020-01-02  102.171269  99.969968  98.682973  100.640755  102.528643
2  2020-01-03  103.171258  99.575237  98.182139  100.574847  101.887811
3  2020-01-04  105.483215  99.308641  97.149381  100.925017  101.490049
4  2020-01-05  107.453175  98.188428  99.575396  101.594411  101.604283
...
360 2020-12-26  92.684784  63.408103  98.288992  117.788079  102.995720
361 2020-12-27  92.688279  62.816639  98.061845  116.605106  102.718260
362 2020-12-28  93.551993  63.597651  96.454800  115.441164  103.566068
363 2020-12-29  93.870037  64.114492  95.747485  113.856107  103.257107
364 2020-12-30  93.855317  64.491011  97.805648  112.640418  102.304226
365 rows x 6 columns

[ ] data.drop_duplicates()

    Unnamed: 0  Stock_1  Stock_2  Stock_3  Stock_4  Stock_5
0  2020-01-01  101.764052  100.160928  99.494642  99.909756  101.761266
1  2020-01-02  102.171269  99.969968  98.682973  100.640755  102.528643
2  2020-01-03  103.171258  99.575237  98.182139  100.574847  101.887811
3  2020-01-04  105.483215  99.308641  97.149381  100.925017  101.490049
4  2020-01-05  107.453175  98.188428  99.575396  101.594411  101.604283
...
360 2020-12-26  92.684784  63.408103  98.288992  117.788079  102.995720
361 2020-12-27  92.688279  62.816639  98.061845  116.605106  102.718260
362 2020-12-28  93.551993  63.597651  96.454800  115.441164  103.566068
363 2020-12-29  93.870037  64.114492  95.747485  113.856107  103.257107
364 2020-12-30  93.855317  64.491011  97.805648  112.640418  102.304226
    
```

## 7. Data Preprocessing

1. Handling missing values, duplicates, and outliers
2. Feature scaling and encoding (if necessary)
3. Data normalization for time series analysis
4. Visual checks for stationarity and trend patterns

```
Untitled4.ipynb
File Edit View Insert Runtime Tools Help
365 rows x 6 columns

[ ] data.columns
Index(['Unnamed: 0', 'Stock_1', 'Stock_2', 'Stock_3', 'Stock_4', 'Stock_5'], dtype='object')

[ ] data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Unnamed: 0   365 non-null    object
1   Stock_1      365 non-null    float64
2   Stock_2      365 non-null    float64
3   Stock_3      365 non-null    float64
4   Stock_4      365 non-null    float64
5   Stock_5      365 non-null    float64
dtypes: float64(5), object(1)
memory usage: 17.2+ KB

[ ] data.isnull().sum()
Unnamed: 0    0
Stock_1       0
Stock_2       0
Stock_3       0
Stock_4       0
Stock_5       0
dtype: int64


[ ] data
   Unnamed: 0  Stock_1  Stock_2  Stock_3  Stock_4  Stock_5
0  2020-01-01  101.764052  100.160928  99.494642  99.909756  101.761266
1  2020-01-02  102.171269  99.969968  98.682973  100.640755  102.528643
2  2020-01-03  103.171258  99.575237  98.182139  100.574847  101.887811
3  2020-01-04  105.483215  99.308641  97.149381  100.925017  101.490049
4  2020-01-05  107.453175  98.188428  99.575396  101.594411  101.604283
...
360 2020-12-26  92.684784  63.408103  98.288992  117.788079  102.995720
361 2020-12-27  92.688279  62.816639  98.061845  116.605106  102.718260
362 2020-12-28  93.551993  63.597651  96.454800  115.441164  103.566068
363 2020-12-29  93.870037  64.114492  95.747485  113.856107  103.257107
364 2020-12-30  93.855317  64.491011  97.805648  112.640418  102.304226
365 rows x 6 columns

[ ] data.drop_duplicates(inplace=True)

[ ] data
   Unnamed: 0  Stock_1  Stock_2  Stock_3  Stock_4  Stock_5
0  2020-01-01  101.764052  100.160928  99.494642  99.909756  101.761266
1  2020-01-02  102.171269  99.969968  98.682973  100.640755  102.528643
2  2020-01-03  103.171258  99.575237  98.182139  100.574847  101.887811
3  2020-01-04  105.483215  99.308641  97.149381  100.925017  101.490049
```




## 8. Exploratory Data Analysis (EDA)

1. Data visualization (e.g., line plots, heatmaps, correlation matrices)
2. Identifying trends, seasonality, and patterns in the data
3. Key takeaways and insights from the EDA phase



## 9. Feature Engineering

1. Creating lag features, moving averages, or technical indicators
2. Dimensionality reduction if required
3. Justifying feature selection for better model performance



AI-Driven Stock Price Prediction.ipynb


+ < > + T
Connect

↑ ↓ ↻ ↗ ↘ ⋮

### Implementation file

stock-price-prediction/ |— data/ | — AAPL.csv ←  
 downloaded dataset |— model/ | — lstm\_stock\_model.h5  
 ← saved trained model |— stock\_predictor.ipynb ← Jupyter  
 notebook for training |— app.py ← Streamlit web app  
 (optional)

```
[ ] import pandas as pd

# Load your CSV file
df = pd.read_csv("Apple Dataset.csv") # Update filename
print(df.head())
print(df.columns)
```

	Date	Open	High	Low	
0	1980-12-12	0.128348	0.128906	0.128348	0
1	1980-12-15	0.122210	0.122210	0.121652	0
2	1980-12-16	0.113281	0.113281	0.112723	0
3	1980-12-17	0.115513	0.116071	0.115513	0
4	1980-12-18	0.118862	0.119420	0.118862	0

Index(['Date', 'Open', 'High', 'Low', 'Close',

### Preprocessing & Normalization

```
[ ] import numpy as np
from sklearn.preprocessing import MinMaxScaler

# Use only the 'Close' column
data = df[['Close']].values

# Normalize
scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(data)

# Create sequences
def create_sequences(data, sequence_length):
    X, y = [], []
    for i in range(sequence_length, len(data)):
        X.append(data[i-sequence_length:i])
        y.append(data[i])
    return np.array(X), np.array(y)

sequence_length = 60
X, y = create_sequences(scaled_data, sequence_length)

# Train-test split
split = int(0.8 * len(X))
```



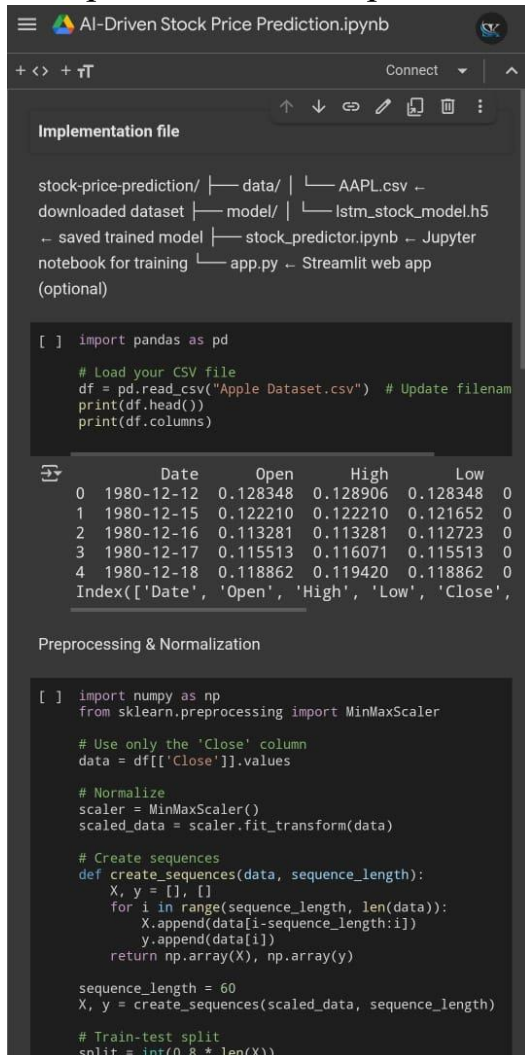
## 10. Model Building

Baseline models (e.g., ARIMA, SARIMA)

Advanced models (e.g., LSTM, GRU, Prophet)

Model training and hyperparameter tuning

Comparison of model performance and selection of the best approach



The screenshot shows a Jupyter Notebook interface with the title "AI-Driven Stock Price Prediction.ipynb". The notebook is divided into sections: "Implementation file" and "Preprocessing & Normalization".

**Implementation file**

stock-price-prediction/ |— data/ | |— AAPL.csv —  
downloaded dataset |— model/ | |— lstm\_stock\_model.h5  
— saved trained model |— stock\_predictor.ipynb — Jupyter  
notebook for training |— app.py — Streamlit web app  
(optional)

```
[ ] import pandas as pd

# Load your CSV file
df = pd.read_csv("Apple Dataset.csv") # Update filename
print(df.head())
print(df.columns)
```

	Date	Open	High	Low	
0	1980-12-12	0.128348	0.128906	0.128348	0
1	1980-12-15	0.122210	0.122210	0.121652	0
2	1980-12-16	0.113281	0.113281	0.112723	0
3	1980-12-17	0.115513	0.116071	0.115513	0
4	1980-12-18	0.118862	0.119420	0.118862	0

Index(['Date', 'Open', 'High', 'Low', 'Close'], dtype='object')

**Preprocessing & Normalization**

```
[ ] import numpy as np
from sklearn.preprocessing import MinMaxScaler

# Use only the 'Close' column
data = df[['Close']].values

# Normalize
scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(data)

# Create sequences
def create_sequences(data, sequence_length):
    X, y = [], []
    for i in range(sequence_length, len(data)):
        X.append(data[i-sequence_length:i])
        y.append(data[i])
    return np.array(X), np.array(y)

sequence_length = 60
X, y = create_sequences(scaled_data, sequence_length)

# Train-test split
split = int(0.8 * len(X))
```

## 11. Model Evaluation

Evaluation metrics (e.g., MAE, RMSE, MAPE)

Visual evaluation (e.g., prediction vs. actual plots)

Error analysis to improve model robustness

```
AI-Driven Stock Price Prediction.ipynb

[ ] sequence_length = 60
X, y = create_sequences(scaled_data, sequence_length)

# Train-test split
split = int(0.8 * len(X))
X_train, X_test = X[:split], X[split:]
y_train, y_test = y[:split], y[split:]

# Reshape for LSTM
X_train = X_train.reshape((X_train.shape[0], X_train.sh
X_test = X_test.reshape((X_test.shape[0], X_test.shape[

Build and Train LSTM Model

[ ] from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropou

model = Sequential([
    LSTM(50, return_sequences=True, input_shape=(X_train
    Dropout(0.2),
    LSTM(50),
    Dropout(0.2),
    Dense(1)
])

model.compile(optimizer='adam', loss='mean_squared_erro
model.fit(X_train, y_train, epochs=20, batch_size=32, v

Epoch 1/20
/usr/local/lib/python3.11/dist-packages/keras/
super().__init__(**kwargs)
273/273 11s 34ms/step
Epoch 2/20
273/273 9s 32ms/step
Epoch 3/20
273/273 9s 33ms/step
Epoch 4/20
273/273 9s 34ms/step
Epoch 5/20
273/273 9s 30ms/step
Epoch 6/20
273/273 10s 30ms/step
Epoch 7/20
273/273 11s 32ms/step
Epoch 8/20
273/273 10s 32ms/step
Epoch 9/20
273/273 10s 38ms/step
Epoch 10/20
273/273 9s 32ms/step
Epoch 11/20
```

## 12. Deployment

Deploy on platforms like Streamlit, Gradio, or Flask

Include a public link and sample output

User interface and interaction design

## 13. Source code

```
# Importing necessary libraries

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import MinMaxScaler

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import LSTM, Dense, Dropout

from sklearn.metrics import mean_squared_error,

mean_absolute_error

import yfinance as yf
```

## # 1. Data Collection

```
stock_symbol = 'AAPL' # Apple Inc.  
  
data = yf.download(stock_symbol, start='2015-01-01',  
end='2025-01-01')  
  
data.reset_index(inplace=True)
```

## # 2. Data Preprocessing

```
# Extracting the 'Close' price for modeling  
  
prices = data[['Date', 'Close']]  
  
prices.set_index('Date', inplace=True)
```

## # Scaling the data

```
scaler = MinMaxScaler()  
  
scaled_data = scaler.fit_transform(prices)
```

## # 3. Feature Engineering - Creating sequences for LSTM

```
sequence_length = 60 # Using 60 days of data to predict the  
next day  
  
X, y = [], []
```

```
for i in range(sequence_length, len(scaled_data)):
```

```
    X.append(scaled_data[i-sequence_length:i])
```

```
    y.append(scaled_data[i])
```

```
X, y = np.array(X), np.array(y)
```

```
# Splitting the data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
```

```
test_size=0.2, random_state=42)
```

```
# 4. Model Building
```

```
model = Sequential()
```

```
model.add(LSTM(50, return_sequences=True,
```

```
input_shape=(X_train.shape[1], X_train.shape[2])))
```

```
model.add(Dropout(0.2))
```

```
model.add(LSTM(50, return_sequences=False))
```

```
model.add(Dropout(0.2))
```

```
model.add(Dense(1))
```

```
model.compile(optimizer='adam',  
loss='mean_squared_error')
```

# 5. Model Training

```
history = model.fit(X_train, y_train, epochs=20,  
batch_size=32, validation_data=(X_test, y_test))
```

# 6. Model Evaluation

```
predictions = model.predict(X_test)  
predictions = scaler.inverse_transform(predictions)  
y_test = scaler.inverse_transform(y_test)
```

```
mse = mean_squared_error(y_test, predictions)  
mae = mean_absolute_error(y_test, predictions)  
print(f"Mean Squared Error (MSE): {mse}")  
print(f"Mean Absolute Error (MAE): {mae}")
```

# Plotting the results

```
plt.figure(figsize=(14,8))
```

```
plt.plot(y_test, color='blue', label='Actual Stock Price')  
  
plt.plot(predictions, color='red', label='Predicted Stock  
Price')  
  
plt.title(f'{stock_symbol} Stock Price Prediction')  
  
plt.xlabel('Days')  
  
plt.ylabel('Stock Price (USD)')  
  
plt.legend()  
  
plt.show()
```

## 14. Future scope

- Incorporate sentiment analysis for market prediction
- Use reinforcement learning for dynamic strategy adjustment
- Explore real-time data pipelines for live market insights

## 15. Team Members and Roles

NAME	ROLE	RESPONSIBLE
Sasivardhan S K	Leader	Data Collection, Data Preprocessing
Joshua Prince S	Member	Feature Engineering
Rubesh Kumar S	Member	Exploratory Data Analysis (EDA),
Mohit Sai Reddy	Member	Model Building, Model Evaluation